

# AWS Solutions Architect Interview Questions & Answers

## Networking & VPC

**Q1: Explain the difference between Security Groups and Network ACLs. When would you use each?**

**Answer:**

### Security Groups:

- Stateful (return traffic automatically allowed)
- Operate at instance/ENI level
- Support allow rules only
- All rules evaluated before decision
- Can reference other security groups

### Network ACLs:

- Stateless (must explicitly allow return traffic)
- Operate at subnet level
- Support both allow and deny rules
- Rules processed in order (lowest number first)
- Cannot reference security groups

### When to use:

- Use Security Groups as primary defense for instance-level security
- Use NACLs for subnet-level protection and to block specific IP ranges
- Example: Block a specific malicious IP at NACL level, allow application traffic at Security Group level

---

**Q2: Design a highly available and secure VPC architecture for a 3-tier web application.**

**Answer:**

## Architecture Components:

### 1. VPC Configuration:

- CIDR: 10.0.0.0/16
- 3 Availability Zones
- Public and Private subnets in each AZ

### 2. Subnet Design:

#### Public Subnets (per AZ):

- 10.0.1.0/24, 10.0.2.0/24, 10.0.3.0/24
- For ALB, NAT Gateways, Bastion hosts

#### Private App Subnets (per AZ):

- 10.0.11.0/24, 10.0.12.0/24, 10.0.13.0/24
- For application servers

#### Private DB Subnets (per AZ):

- 10.0.21.0/24, 10.0.22.0/24, 10.0.23.0/24
- For RDS instances

### 3. Routing:

- Internet Gateway for public subnets
- NAT Gateway in each AZ for private subnet internet access
- Route tables: public (IGW), private-app (NAT), private-db (no internet)

### 4. Security:

Web Tier SG: Allow 443 from 0.0.0.0/0

App Tier SG: Allow 8080 from Web Tier SG only

DB Tier SG: Allow 3306 from App Tier SG only

### 5. High Availability:

- Multi-AZ ALB
- Auto Scaling Groups in multiple AZs
- RDS Multi-AZ deployment
- Multi-AZ NAT Gateways

## Key Design Decisions:

- Separate subnets per tier for security isolation
- Multi-AZ for high availability
- Security groups reference other security groups for maintainability
- No direct internet access to private subnets

### **Q3: How would you connect multiple VPCs across different regions and accounts?**

**Answer:**

**Solution: Transit Gateway with RAM (Resource Access Manager)**

**Architecture:**

1. Hub-and-Spoke Model:

- Central Transit Gateway in primary region
- Peering connections to Transit Gateways in other regions
- Spoke VPCs attach to regional Transit Gateways

2. Implementation Steps:

- a) Create Transit Gateway in each region
- b) Share Transit Gateway using AWS RAM to other accounts
- c) Accept RAM invitation in target accounts
- d) Attach VPCs to Transit Gateway
- e) Configure route tables for inter-VPC communication
- f) Create Transit Gateway Peering between regions

3. Benefits:

- Transitive routing support
- Centralized network management
- Scales to thousands of VPCs
- Reduced complexity vs mesh VPC peering

**Alternative for Service Access:**

- AWS PrivateLink for specific service exposure
- VPC Peering for simple 1-to-1 connections

**Cost Considerations:**

- Transit Gateway: \$0.05/hour + \$0.02/GB data transfer
- VPC Peering: \$0.01/GB (same region)
- Choose based on number of VPCs and data volume

---

**Compute Services**

### **Q4: When would you choose EC2, ECS, EKS, Lambda, or Fargate? Explain your decision criteria.**

**Answer:**

## **EC2 (Elastic Compute Cloud):**

- **Use When:** Full OS control needed, Windows workloads, lift-and-shift migrations
- **Example:** Legacy .NET applications, specific compliance requirements
- **Pros:** Maximum flexibility, any OS/software
- **Cons:** Most operational overhead, patch management

## **ECS (Elastic Container Service):**

- **Use When:** Docker containers, AWS-native solution, simpler than K8s
- **Example:** Microservices architecture without Kubernetes complexity
- **Pros:** Deep AWS integration, simpler than EKS, less overhead
- **Cons:** AWS-specific, limited portability

## **EKS (Elastic Kubernetes Service):**

- **Use When:** Need Kubernetes, multi-cloud strategy, complex orchestration
- **Example:** Organizations with existing K8s expertise, hybrid cloud
- **Pros:** Industry standard, portable, rich ecosystem
- **Cons:** Complex, higher costs, steeper learning curve

## **Lambda:**

- **Use When:** Event-driven, <15 min execution, variable workload
- **Example:** API backends, data processing, scheduled tasks
- **Pros:** No server management, auto-scaling, pay-per-use
- **Cons:** Cold starts, 15-min limit, stateless

## **Fargate:**

- **Use When:** Containers without server management (ECS/EKS)
- **Example:** Microservices with unpredictable traffic
- **Pros:** Serverless containers, no capacity planning
- **Cons:** Higher cost than EC2, less control

## **Decision Tree:**

Stateless, event-driven, <15 min? → Lambda

Containers + no server management? → Fargate

Containers + Kubernetes needed? → EKS

Containers + simpler orchestration? → ECS

Full control needed? → EC2

---

## **Q5: Design an auto-scaling solution for a web application with predictable daily traffic patterns.**

**Answer:**

**Solution: Scheduled + Dynamic Scaling**

## 1. Traffic Pattern Analysis:

- Baseline: 10 instances (midnight-6am)
- Morning ramp: 10→30 instances (6am-9am)
- Peak: 30-50 instances (9am-6pm)
- Evening ramp down: 50→10 instances (6pm-midnight)

## 2. Auto Scaling Configuration:

### a) Scheduled Scaling Actions:

- 6:00 AM: Set min=20, desired=30, max=50
- 9:00 AM: Set min=30, desired=30, max=80
- 6:00 PM: Set min=20, desired=25, max=40
- 11:00 PM: Set min=10, desired=10, max=30

### b) Dynamic Scaling Policies:

#### Target Tracking Policy:

- Metric: CPU Utilization 70%
- Scale-out: Add 20% capacity when >70% for 2 min
- Scale-in: Remove 10% capacity when <50% for 10 min

#### Step Scaling (for spikes):

- 70-80% CPU: +2 instances
- 80-90% CPU: +4 instances
- >90% CPU: +8 instances

## 3. Additional Configuration:

- Warm-up time: 300 seconds
- Health check grace period: 180 seconds
- Scale-in protection during deployments
- Multiple AZs for availability
- ALB with slow start mode (30 seconds)

## 4. Cost Optimization:

- Use Reserved Instances for baseline (10 instances)
- On-Demand for predictable scaling (10-30 instances)
- Spot Instances for burst capacity (above 30)

## Monitoring:

- CloudWatch dashboards for scaling activities
- Alarms for scale-in/out events
- Lambda for notifications

**Q6: How do you handle sudden, unpredictable traffic spikes cost-effectively?**

**Answer:**

**Multi-Layer Defense Strategy:**

## 1. Edge Layer - CloudFront:

- Cache static content (images, CSS, JS)
- Reduce origin load by 60-80%
- DDoS protection with AWS Shield
- Lambda@Edge for simple request routing

## 2. Application Layer - Multiple Options:

### Option A: Lambda Architecture

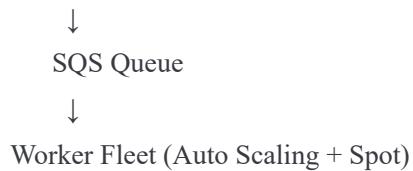
- API Gateway + Lambda
- Automatic scaling to handle spikes
- Pay only for actual requests
- Best for: APIs, simple logic

### Option B: Container + Spot Instances

- ECS/EKS with Fargate Spot
- Baseline on On-Demand
- Burst capacity on Spot (70% cost savings)
- Cluster Autoscaler for EKS

## 3. Queue-Based Decoupling:

Internet → ALB → API (small fleet)



### Benefits:

- Prevent application overload
- Process requests asynchronously
- Queue depth metric triggers scaling

## 4. Database Layer:

- Read Replicas for read-heavy workloads
- ElastiCache (Redis) for frequently accessed data
- DynamoDB with on-demand capacity
- Connection pooling (RDS Proxy)

## 5. Cost Optimization Tactics:

- Spot Instances for burst capacity
- Mixed instance types in Auto Scaling
- Savings Plans for baseline
- Right-sizing based on actual usage

## 6. Implementation Example:

Instance Mix for Web Tier:

- 20% Reserved (baseline): c5.large
- 10% On-Demand (buffer): c5.large
- 70% Spot (burst): c5.large, c5.xlarge, c6i.large

**Real-World Scenario:** Normal load: 100 req/sec → 10 instances Sudden spike: 5,000 req/sec → Queue buffers, Auto Scaling adds 90 instances over 5 minutes, Spot provides 70% of added capacity

---

## Storage & Databases

**Q7: Explain different S3 storage classes and design a lifecycle policy for a data lake.**

**Answer:**

**S3 Storage Classes:**

## 1. S3 Standard:

- Use: Frequently accessed data
- Latency: Milliseconds
- Cost: \$0.023/GB/month
- Availability: 99.99%

## 2. S3 Intelligent-Tiering:

- Use: Unknown/changing access patterns
- Auto-moves between access tiers
- Cost: \$0.023-\$0.0125/GB + monitoring fee

## 3. S3 Standard-IA:

- Use: Infrequent access (monthly)
- Cost: \$0.0125/GB + retrieval fee
- Min storage: 30 days

## 4. S3 One Zone-IA:

- Use: Recreational infrequent data
- Cost: \$0.01/GB (20% cheaper than Standard-IA)
- Single AZ (99.5% availability)

## 5. S3 Glacier Instant Retrieval:

- Use: Archive with instant access (quarterly)
- Cost: \$0.004/GB
- Retrieval: Milliseconds

## 6. S3 Glacier Flexible Retrieval:

- Use: Archive (1-2 times/year)
- Cost: \$0.0036/GB
- Retrieval: Minutes to hours

## 7. S3 Glacier Deep Archive:

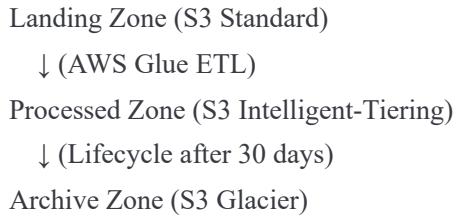
- Use: Long-term compliance (7-10 years)
- Cost: \$0.00099/GB
- Retrieval: 12-48 hours

## Data Lake Lifecycle Policy:

json

```
{
  "Rules": [
    {
      "Id": "Raw-Data-Lifecycle",
      "Filter": {"Prefix": "raw-data/"},
      "Status": "Enabled",
      "Transitions": [
        {
          "Days": 30,
          "StorageClass": "STANDARD_IA"
        },
        {
          "Days": 90,
          "StorageClass": "GLACIER_IR"
        },
        {
          "Days": 365,
          "StorageClass": "DEEP_ARCHIVE"
        }
      ]
    },
    {
      "Id": "Processed-Data-Lifecycle",
      "Filter": {"Prefix": "processed/"},
      "Status": "Enabled",
      "Transitions": [
        {
          "Days": 7,
          "StorageClass": "INTELLIGENT_TIERING"
        }
      ]
    },
    {
      "Id": "Temp-Data-Expiration",
      "Filter": {"Prefix": "temp/"},
      "Status": "Enabled",
      "Expiration": {
        "Days": 7
      }
    }
  ]
}
```

## Data Lake Architecture:



## Q8: How do you choose between RDS, Aurora, DynamoDB, and Redshift?

**Answer:**

**Decision Matrix:**

**Amazon RDS:**

- **Use For:** Traditional relational databases, lift-and-shift
- **Engines:** MySQL, PostgreSQL, Oracle, SQL Server, MariaDB
- **Max Size:** 64 TB
- **Performance:** Up to 16,000 IOPS (io1)
- **Best For:** Existing applications, ACID compliance, complex queries
- **Example:** E-commerce order management, CRM systems

**Amazon Aurora:**

- **Use For:** High-performance relational database
- **Compatible:** MySQL, PostgreSQL
- **Performance:** 5x MySQL, 3x PostgreSQL performance
- **Max Size:** 128 TB (auto-scaling storage)
- **Read Replicas:** Up to 15
- **Best For:** High-throughput OLTP, need for read scalability
- **Global Database:** <1 sec cross-region replication
- **Example:** Gaming leaderboards, financial transactions

**Amazon DynamoDB:**

- **Use For:** NoSQL key-value/document store
- **Performance:** Single-digit millisecond latency at any scale
- **Scaling:** Automatic, handles millions of requests/sec

- **Data Model:** Key-value, document
- **Best For:** Mobile/web apps, gaming, IoT, real-time
- **Global Tables:** Multi-region, active-active
- **Example:** User profiles, shopping carts, session management

### Amazon Redshift:

- **Use For:** Data warehouse, analytics, OLAP
- **Performance:** Columnar storage, MPP architecture
- **Max Size:** Petabyte-scale
- **Best For:** Business intelligence, complex analytical queries
- **Not For:** OLTP, real-time queries
- **Example:** Sales analytics, data marts, reporting

### Selection Flow:

```

Need relational + SQL?
├─ Yes → OLTP or OLAP?
  |  ├─ OLTP → High performance needed?
  |  |  ├─ Yes → Aurora
  |  |  └─ No → RDS
  |  └─ OLAP → Redshift
└─ No → NoSQL → DynamoDB

```

### Hybrid Example:

```

User Profile → DynamoDB (fast K-V lookup)
Transactions → Aurora (ACID, relational)
Analytics → Redshift (aggregations, reporting)
Session Data → ElastiCache (in-memory)

```

---

## Q9: Design a database migration strategy from on-premises Oracle to AWS with minimal downtime.

### Answer:

**Migration Strategy: AWS DMS with CDC**

**Phase 1: Assessment (1-2 weeks)**

### 1. Database Analysis:

- Size: 5 TB
- Transaction rate: 10,000 TPS
- Dependencies: 50 applications
- Downtime tolerance: <1 hour

### 2. Target Selection:

- Option A: RDS for Oracle (license-included)  
Option B: Aurora PostgreSQL (cost-effective)  
Option C: DynamoDB (if NoSQL suitable)

Decision: Aurora PostgreSQL (70% cost savings)

### 3. Schema Conversion Tool (SCT):

- Assess compatibility
- Generate conversion report
- Identify manual changes needed

## Phase 2: Setup (1 week)

### 1. Network Connectivity:

- VPN or Direct Connect to AWS
- Security groups: Allow Oracle port from DMS

### 2. Create Target Database:

- Aurora PostgreSQL cluster
- Multi-AZ for HA
- Appropriate instance size (r6g.4xlarge)

### 3. DMS Configuration:

- Replication instance: dms.c5.4xlarge
- Source endpoint: On-premises Oracle
- Target endpoint: Aurora PostgreSQL

## Phase 3: Migration (2-3 weeks)

### 1. Full Load Phase:

- DMS performs initial bulk copy
- Estimated time: 72 hours for 5 TB
- Run during low-traffic period

### 2. Change Data Capture (CDC):

- Captures ongoing changes during migration
- Keeps source and target in sync
- Monitor replication lag

### 3. Testing:

- Parallel run: Applications test against Aurora
- Performance testing
- Data validation (row counts, checksums)
- Application compatibility testing

## Phase 4: Cutover (1 hour)

### Cutover Steps:

1. Enable read-only mode on source (T+0 min)
2. Wait for CDC lag to reach zero (T+5 min)
3. Validate data consistency (T+10 min)
4. Update application connection strings (T+15 min)
5. Perform smoke tests (T+20 min)
6. Enable write traffic to Aurora (T+30 min)
7. Monitor for issues (T+30-60 min)

### Rollback Plan:

- Keep source database available for 48 hours
- Can switch back connection strings if needed

## Phase 5: Post-Migration

1. Monitoring:

- Database performance metrics
- Application error rates
- Query performance

2. Optimization:

- Create indexes based on workload
- Adjust Aurora parameters
- Enable Performance Insights

3. Decommission:

- After 2 weeks of stable operation
- Archive source database

### **Alternative: Zero-Downtime with Bidirectional Replication**

- Use Oracle GoldenGate for bi-directional replication
  - Allows gradual application cutover
  - Higher complexity and cost
- 

### **Q10: How do you design a globally distributed, low-latency database architecture?**

**Answer:**

**Solution: Multi-Region Active-Active Architecture**

**Option 1: DynamoDB Global Tables**

#### Architecture:

- DynamoDB tables in 5 regions (us-east-1, eu-west-1, ap-southeast-1, etc.)
- Multi-region, multi-active replication
- Last-writer-wins conflict resolution

#### Implementation:

1. Create base table in primary region
2. Enable streams
3. Add replica regions via Global Tables
4. Configure Route 53 geolocation routing

#### Benefits:

- Single-digit millisecond latency globally
- <1 second replication between regions
- 99.999% availability SLA
- Automatic conflict resolution

#### User Flow:

User (Tokyo) → Route 53 → API Gateway (ap-northeast-1)  
→ Lambda → DynamoDB (ap-northeast-1)  
→ Async replication to other regions

## Option 2: Aurora Global Database

#### Architecture:

- Primary region: Read-write (us-east-1)
- Secondary regions: Read-only (eu-west-1, ap-southeast-1)
- Physical replication (<1 second lag)

#### Implementation:

1. Create Aurora cluster in primary region
2. Add global database cluster
3. Add secondary regions with read replicas
4. Configure Route 53 latency-based routing

#### Routing Logic:

- Writes: Always route to primary region
- Reads: Route to nearest region

#### Disaster Recovery:

- RTO: <1 minute (promote secondary)
- RPO: <1 second

#### Read Flow:

User (London) → Route 53 (latency routing)  
→ ALB (eu-west-1)  
→ App Servers  
→ Aurora read replica (eu-west-1)

#### Write Flow:

User (London) → Route 53  
→ ALB (us-east-1)  
→ App Servers (us-east-1)  
→ Aurora primary (us-east-1)  
→ Replicates to eu-west-1

## Option 3: Hybrid with Caching

Architecture:

```
Internet
  ↓
CloudFront (edge locations worldwide)
  ↓
Route 53 (geolocation routing)
  ↓
Regional API Gateways
  ↓
ElastiCache Global Datastore (Redis)
  ↓
Aurora Global Database
```

Caching Strategy:

- User profiles: Cache 1 hour
- Product catalog: Cache 5 minutes
- Inventory: No cache (real-time)
- Session data: Redis with TTL

Benefits:

- <10ms latency for cached data
- 90% cache hit rate reduces DB load
- Global consistency for writes

## Comparison:

Feature	DynamoDB Global	Aurora Global	Hybrid
Write Latency	<10ms (local)	<100ms (cross-region)	<10ms (cached)
Read Latency	<10ms	<10ms (local replica)	<5ms (cached)
Conflict Resolution	Automatic	Manual	Application-level
Cost	\$\$\$	\$\$	\$\$\$\$
Complexity	Low	Medium	High

## Recommendation:

- **High write throughput, simple K-V:** DynamoDB Global Tables
- **Complex relational queries:** Aurora Global Database + ElastiCache
- **Read-heavy, low latency:** Hybrid with aggressive caching

## **Security & Compliance**

**Q11: Explain your approach to implementing defense-in-depth security for AWS resources.**

**Answer:**

**Multi-Layer Security Strategy:**

## Layer 1: Edge Protection

- CloudFront with AWS WAF
- AWS Shield (DDoS protection)
- Geo-blocking for restricted regions
- Rate limiting

## Layer 2: Network Security

- VPC with private subnets
- NACLs for subnet-level filtering
- Security Groups (stateful firewall)
- VPC Flow Logs for monitoring
- AWS Network Firewall for deep packet inspection

## Layer 3: Identity & Access

- IAM roles with least privilege
- MFA enforcement
- AWS SSO for centralized access
- Service Control Policies (SCPs)
- Regular access reviews

## Layer 4: Application Security

- ALB with WAF rules (OWASP Top 10)
- API Gateway with throttling
- Secrets Manager for credentials
- Certificate Manager for TLS
- Inspector for vulnerability scanning

## Layer 5: Data Security

- Encryption at rest (KMS)
- Encryption in transit (TLS 1.2+)
- S3 bucket policies
- RDS encryption
- Macie for sensitive data discovery

## Layer 6: Monitoring & Response

- CloudTrail (API logging)
- GuardDuty (threat detection)
- Security Hub (centralized view)
- Config (compliance monitoring)
- EventBridge + Lambda for automated response

## **Implementation Example:**

## Web Application Security Stack:

### 1. Perimeter:

CloudFront → WAF Rules:

- SQL injection protection
- XSS filtering
- Bot control
- Rate limiting: 2000 req/5min per IP

### 2. Network:

Public Subnet SG:

- Allow 443 from CloudFront IPs only

Private Subnet SG:

- Allow 8080 from ALB SG only

Database SG:

- Allow 3306 from App SG only

### 3. Access Control:

EC2 Instance Role:

- Read from S3 bucket (specific bucket)
- Write to CloudWatch Logs
- Read from Secrets Manager (specific secret)

### 4. Data Protection:

S3 Bucket Policy:

- Deny non-HTTPS access
- Deny public access
- Require encryption
- Versioning enabled

RDS:

- Encryption at rest (KMS CMK)
- Automated backups encrypted
- SSL/TLS enforced

### 5. Monitoring:

CloudWatch Alarms:

- Failed login attempts >5 in 5 min
- Root account usage
- Security group changes
- IAM policy changes

GuardDuty Findings:

- Lambda auto-isolates compromised instances

---

## **Q12: How do you implement encryption for data at rest and in transit across AWS services?**

**Answer:**

**Encryption Strategy:**

**At Rest Encryption:**

#### 1. S3:

- SSE-S3: AWS-managed keys
- SSE-KMS: Customer managed keys (CMK)
- SSE-C: Customer-provided keys

Recommended: SSE-KMS with CMK

```
{  
  "Version": "2012-10-17",  
  "Statement": [  
    {  
      "Effect": "Deny",  
      "Principal": "*",  
      "Action": "s3:PutObject",  
      "Resource": "arn:aws:s3:::my-bucket/*",  
      "Condition": {  
        "StringNotEquals": {  
          "s3:x-amz-server-side-encryption": "aws:kms"  
        }  
      }  
    }  
  ]  
}
```

#### 2. EBS:

- Enable encryption by default (regional setting)
- Uses KMS CMK
- Snapshots automatically encrypted

AWS CLI:

```
aws ec2 enable-ebs-encryption-by-default --region us-east-1
```

#### 3. RDS/Aurora:

- Enable at creation (cannot add later)
- Uses KMS CMK
- Automated backups, snapshots encrypted
- Read replicas must use same encryption

#### 4. DynamoDB:

- Encryption enabled by default
- AWS owned, AWS managed, or Customer managed KMS keys
- No performance impact

#### 5. EFS:

- Enable encryption at creation
- Uses KMS CMK
- Applies to data and metadata

## In Transit Encryption:

### 1. Application Load Balancer:

- HTTPS listener with ACM certificate
- SSL/TLS termination at ALB
- Backend encryption optional

Configuration:

- Protocol: HTTPS
- Security Policy: ELBSecurityPolicy-TLS-1-2-2017-01
- Certificate: ACM or imported

### 2. API Gateway:

- Enforces HTTPS by default
- Custom domain with ACM certificate
- Regional or edge-optimized endpoints

### 3. RDS:

- Force SSL connections

MySQL:

```
REQUIRE SSL;
```

PostgreSQL:

```
rds.force_ssl = 1
```

### 4. S3:

- Bucket policy to deny non-HTTPS

```
{  
  "Effect": "Deny",  
  "Principal": "*",  
  "Action": "s3:*",  
  "Resource": "arn:aws:s3:::my-bucket/*",  
  "Condition": {  
    "Bool": {"aws:SecureTransport": "false"}  
  }  
}
```

### 5. VPC:

- VPN with IPsec
- Direct Connect with MACsec
- PrivateLink for service access

## Key Management:

## KMS Best Practices:

### 1. Key Hierarchy:

- Separate CMKs per environment (dev/prod)
- Separate CMKs per data classification
- Automatic key rotation enabled (yearly)

### 2. Key Policy:

```
{  
  "Version": "2012-10-17",  
  "Statement": [  
    {  
      "Sid": "Enable IAM policies",  
      "Effect": "Allow",  
      "Principal": {"AWS": "arn:aws:iam::123456789012:root"},  
      "Action": "kms:*",  
      "Resource": "*"  
    },  
    {  
      "Sid": "Allow S3 to use key",  
      "Effect": "Allow",  
      "Principal": {"Service": "s3.amazonaws.com"},  
      "Action": ["kms:Decrypt", "kms:GenerateDataKey"],  
      "Resource": "*"  
    }  
  ]  
}
```

### 3. Monitoring:

- CloudTrail logs all KMS API calls
- CloudWatch alarms for key usage
- Audit key policies regularly

## High Availability & Disaster Recovery

### Q13: Explain RTO and RPO. Design solutions for different recovery requirements.

#### Answer:

#### Definitions:

- **RTO (Recovery Time Objective):** Maximum acceptable downtime
- **RPO (Recovery Point Objective):** Maximum acceptable data loss

## **DR Strategies (Cost ↔ Recovery Time Trade-off):**

### **1. Backup and Restore (High RTO/RPO - Lowest Cost)**

RTO: Hours to days

RPO: Hours

Cost: \$

Architecture:

- Regular snapshots to S3
- CloudFormation templates for infrastructure
- Automated backups

Use Case: Non-critical applications, cost-sensitive

Implementation:

Primary Region:

- EC2 with daily AMI snapshots
- RDS with automated backups (7-day retention)
- S3 with versioning

DR Steps:

1. Create VPC from CloudFormation (30 min)
2. Launch EC2 from latest AMI (15 min)
3. Restore RDS from snapshot (45 min)
4. Update DNS (5 min)

Total RTO: ~2 hours

### **2. Pilot Light (Medium RTO/RPO - Low Cost)**

RTO: 10 minutes to 1 hour

RPO: Minutes

Cost: \$\$

Architecture:

- Minimal resources always running in DR region
- Database replication active
- Application servers off/minimal

Use Case: Business-critical applications

Implementation:

DR Region (Always Running):

- VPC and subnets configured
- RDS read replica (continuous replication)
- AMIs updated weekly
- Data replication via S3 Cross-Region Replication

Failover Steps:

1. Promote RDS read replica (5 min)
2. Launch EC2 from AMIs via Auto Scaling (5 min)
3. Update Route 53 (2 min)
4. Warm up application (3 min)

Total RTO: 15-20 minutes

RPO: ~5 minutes (replication lag)

### **3. Warm Standby (Low RTO/RPO - Medium Cost)**

RTO: Minutes

RPO: Seconds to minutes

Cost: \$\$\$

Architecture:

- Scaled-down version running in DR region
- Continuous data replication
- Can handle some load immediately

Implementation:

DR Region (Running):

- Auto Scaling: min=2, max=10 (vs prod: min=10, max=50)
- Smaller instance types (t3.medium vs c5.xlarge)
- Aurora Global Database (replication lag <1 sec)
- Route 53 weighted routing (90% primary, 10% DR)

Failover Steps:

1. Increase Auto Scaling capacity (2 min)
2. Update Route 53 weights to 0% primary, 100% DR (1 min)

Total RTO: 5 minutes

RPO: <1 minute

#### 4. Multi-Site Active-Active (Minimal RTO/RPO - Highest Cost)

RTO: Automatic (zero downtime)

RPO: Zero to seconds

Cost: \$\$\$\$

Architecture:

- Full capacity in multiple regions
- Active traffic in all regions
- Real-time data replication

Implementation:

Global Setup:

- DynamoDB Global Tables (multi-region, multi-active)
- Route 53 latency/geo routing
- CloudFront for global edge caching
- Identical infrastructure in all regions

Each Region:

- Full production capacity
- Auto Scaling based on traffic
-

